







A. Interventions

 SYSTEM PROMPT	Rich prompt: "You are [one of 15 historical figures]..."
 IN-CONTEXT LEARNING	32 biographical Q&A pairs prefilled in context (person is never named)
 SUPERVISED FINE-TUNING	LoRA on 300 in-character Q&A pairs per persona
 OPEN-CHARACTER TRAINING	Per-persona constitution → DPO distillation from teacher → introspective SFT
 EMERGENT MISALIGNMENT	Bad-medical-advice fine-tuning. Not one of the 15 historical persons



B. Probe Uplift Scores



C. Black-Box Behaviour




D. Conclusions



Role-play is shallow. It acts like the character, with little shift in what it represents as true.



We observe a spectrum of internalization, with shallow methods on one end and OCT and EM on the other.



EM shifts the model toward internally representing harmful falsehoods as true – beyond its training domain, most strongly on historical atrocities.